RecallPilot

PRODUCT PLAN

Version 1.0

AI RECALL & REVENUE RECOVERY FOR DENTAL PRACTICES

Bring patients back. Prove the revenue.

The product strategy, architecture, delivery plan, pricing model and brand system for building RecallPilot as a real vertical SaaS product.

DATE	OWNER	STATUS
8 November 2026	RecallPilot	Confidential

00 CONTENTS

Everything needed to build the product.

This document consolidates the business case, product scope, technical architecture, delivery plan and brand system for RecallPilot. It is written to be the working plan for the first 12 weeks of product development.

01	Executive summary	Wedge, decisions, success criteria
02	Business context	Product layers, journey, boundaries
03	Pricing and business model	Plans, economics, risks, validation
04	Product plan	MVP scope, workflows, architecture, data
05	Delivery plan	Roadmap, R0–R1, acceptance criteria
06	Brand system	Logos, colors, typography, UI tokens

DOCUMENT	VERSION	DATE	CLASSIFICATION
Product plan	1.0	8 Nov 2026	Confidential

Locked decisions

The product is built in a new, private repository: **v-cut127/RecallPilot-App**. The current repository remains the landing page. The first release is a **free recall audit plus CSV import**. The stack is **Next.js 15 + Supabase EU + pg-boss**, using synthetic data until a real clinic export is available. The UI and report are English-first; the data model is prepared for German and Dutch.

01 EXECUTIVE SUMMARY

The wedge is a number, not a feature list.

RecallPilot is an AI recall and revenue-recovery agent for dental practices in Europe. The product does not sell "AI for dentists." It sells a monthly outcome: **"We recovered 27 appointments and €8,430 of production you would otherwise have lost."**

Build the smallest closed loop that produces that number: **CSV import → free recall audit → email pilot → booking → recovered-revenue report.**

€199–599

Monthly plans, priced against recovered production rather than software seats.

10–12 wks

Pilot-ready MVP with one full-time developer and disciplined scope.

10 pilots

Validation target before investing in deep PMS integrations, voice or WhatsApp.

Why this can win

OUTCOME, NOT TOOLING

A report the practice owner forwards.

Every feature exists to produce a credible recovered-revenue number. The audit is the lead magnet; the monthly report is the retention mechanism.

NARROW FIRST RELEASE

Email-first, CSV-first, human-approved.

No deep PMS write-back, no WhatsApp, no voice in v0.1. The first release proves demand and lift with the lowest technical and compliance risk.

TRUST BY DESIGN

Holdout group from day one.

Attribution is the product. A control group prevents the company from overclaiming and gives the practice a reason to believe the report.

SAFETY AS A SALES ASSET

Clinical questions always escalate.

RecallPilot never diagnoses, promises insurance coverage or gives clinical advice. That boundary is enforced in the agent and explained to every practice.

02 BUSINESS CONTEXT

A three-layer product with a hard boundary.

RecallPilot connects to the practice-management system, decides who to contact and why, and reports the appointments and production recovered. Each layer is necessary, but the outcome layer is the wedge.

LAYER 1

Data

Read/write connection to the PMS: patients, appointments, recall intervals, treatment plans, contact consent. Writes back bookings and conversation outcomes.

CSV first · Open Dental later

LAYER 2

Agent

A rules engine plus LLM decides who to contact, when, on which channel, in which language, with which message. Handles replies, objections and escalation.

Rules engine · Human approval

LAYER 3

Outcome

A dashboard and monthly report attributing booked appointments and recovered production to the agent, with a control group so the practice can trust the number.

Holdout · Monthly report

Product boundary

RecallPilot is **not** a clinical decision-support tool and **not** a full practice-management system. It must never diagnose, promise insurance coverage or give clinical advice. Every ambiguous clinical message is escalated to the practice. This boundary is both a safety requirement and a sales asset.

Pilot guardrails

COHORT

100–200 low-risk patients

Exclude vulnerable patients, open complaints and anyone with ambiguous consent.

MEASUREMENT

90/10 holdout

A control group from the first send so lift is measured, not assumed.

CONTROL

Human approval

Every message is reviewed during the pilot. Auto-send is earned later.

CADENCE

Weekly review

Response, bookings, front-desk load and every objection are reviewed weekly.

02 RECALL JOURNEY

From overdue list to booked appointment.

The customer journey

Step	Action	Detail
1	**Connect the PMS**	OAuth/API if available; otherwise secure CSV upload or read-only export. Validate patient counts, appointment history, recall intervals and consent fields.
2	**Run the free recall audit**	Within 24 hours: how many patients are overdue, estimated annual production at risk, and sample messages in the practice's language.
3	**Configure the agent**	Recall rules, languages, channels, quiet hours, providers and escalation contacts. Approve scripts and the "do not say" list.
4	**Pilot for 30 days**	Start with 100–200 low-risk patients. Measure response, bookings and front-desk time. The practice can pause the agent at any time.
5	**Review the outcome**	Weekly check-in and monthly recovered-revenue report. If the agreed threshold is not met, the practice does not continue or pays a reduced fee.
6	**Expand**	Add treatment follow-up, more locations, voice and review requests. Move to annual billing with a discount.

What the AI automates vs. what requires a human

AI AUTOMATES

- Segmentation of overdue patients and recall rules.
- Message personalization, translation and send-time optimization.
- Reply classification, objection handling and booking into open slots.
- Rescheduling, confirmation and no-show follow-up.
- Reporting and attribution of recovered appointments.
- Opt-out handling and suppression lists.

HUMAN INTERVENTION REQUIRED

- Clinical questions ("Is this bleeding normal?") and any medical advice.
- Insurance/benefit exceptions and payment disputes.
- Complaints, vulnerable patients and ambiguous requests.
- Final approval of scripts, escalation policies and consent rules.
- Manual audit of a sample of conversations for quality.
- Legal/DPO review of the data-processing agreement and DPIA.

Market sequencing

Start with Germany and the Netherlands, the two markets named in the validation plan. Ship the product UI and report in English first, with the data model, templates and locale columns ready for German and Dutch. Add local-language QA before sending in those markets.

03 PRICING AND BUSINESS MODEL

Price against recovered production.

A recovered hygiene appointment is worth €100–250 immediately and €300–800 per year in recurring production. At €349/month, the practice needs roughly 2–4 recovered appointments to break even — a very easy sales conversation.

Tier	Monthly	Features and limits	Target customer	AI/API cost	Gross margin
Starter	**€199**	1 location; 500 patients contacted/month; SMS + email; 1 language; recall dashboard; monthly report; email support.	Solo practice or first-time user; 1,000–2,500 active patients.	€25–45	77–87%
Growth	**€349**	2,000 patients/month; WhatsApp; 2 languages; treatment follow-up; voice add-on (€99 for 300 min); priority support.	1–3 location practice with an active recall problem.	€60–110	68–83%
Scale	**€599**	3 locations; 6,000 patients/month; unlimited languages; 300 voice minutes included; API; dedicated onboarding; quarterly business review.	Multi-location group or dental service organisation.	€120–220	63–80%

Cost assumptions: email near-zero; SMS €0.04–0.08 per message; WhatsApp €0.03–0.08 per conversation; LLM €0.01–0.05 per conversation; voice €0.10–0.25 per minute; infrastructure €5–30 per account. An SMS-heavy mix can reduce gross margin by 10–20 points, so plans must include message credits or overage pricing.

Pricing logic

Price against the value of a recovered appointment, not against other software. Use a 30-day pilot with a clear success threshold. After attribution is trusted, introduce a performance fee: €12 per recovered appointment after the first 10, capped at €299/month.

Add-ons and expansion revenue

Add-on	Price	What it includes
Voice	€99/month	300 minutes. Higher margin than SMS; useful for older patient demographics.
Treatment follow-up	+€99/month	Higher value because unscheduled treatment is worth more than hygiene.
Review requests	+€49/month	Sends a review request after a completed appointment.
Performance fee	€12/appointment	After the first 10 recovered appointments, capped at €299/month.
Setup	€299 one-time	Waived on annual billing. Covers CSV cleanup, script approval and staff training.

03 UNIT ECONOMICS

Churn is the variable to watch.

Revenue scenarios

Customers	Average monthly price	MRR	ARR	Comment
10	€199	€1,990	€23,880	Design partners; founder-led sales; pre-scale.
50	€299	€14,950	€179,400	First repeatable outbound motion; one CS hire.
100	€349	€34,900	€418,800	Not profitable with a 4-person team; need ~160–200 customers.
250	€349	€87,250	€1,047,000	Profitable with a lean team; partner channel matters.
500	€349	€174,500	€2,094,000	Requires sales capacity, onboarding and a second product line.

Unit economics

Blended CAC	€600
ARPU	€299
Gross margin	75%
Monthly churn	3.5%
LTV (3.5% churn)	€6,407
LTV / CAC	10.7×
CAC payback	2.7 months
LTV at 5% churn	€4,485

THE VARIABLE TO WATCH

Protect the first 90 days.

LTV remains attractive at 5% churn, but the first 90 days are fragile. The product must deliver visible recovered revenue quickly, or the practice will read the tool as another inbox and cancel.

Keep the team lean, raise ARPU through treatment follow-up and multi-location plans, or reach 200+ customers before hiring aggressively.

Margin guardrails

Email is near-zero; SMS costs are controlled with credits; WhatsApp template fees are reviewed monthly; voice is an add-on, not a default. Re-price if the channel mix shifts toward SMS or voice.

03 CHURN RISKS

Retention is the growth strategy.

A low-ACV SaaS is not profitable at 100 customers with a full team. Retention depends on visible recovered revenue within the first 90 days and a front desk that does not feel replaced.

Churn risks and mitigations

Risk	Why it happens	Mitigation
Practice does not see recovered revenue	Attribution is weak or patient response rate is low.	Control group, booking-source tracking, weekly reports, 30-day pilot with a clear success threshold.
Front desk rejects the tool	Fear of losing control or extra work.	Human-in-the-loop inbox, staff training, “assistant not replacement” positioning, quick wins in week one.
PMS integration breaks	API changes, export errors, vendor restrictions.	Monitoring, CSV fallback, integration status page, partner agreements, error alerts.
Patient complaints or opt-outs	Too many messages, wrong tone, language errors.	Frequency caps, quiet hours, easy opt-out, human review of scripts, local-language QA.
Practice closure or merger	Market consolidation.	Multi-location plans, diversify across countries, expand to veterinary/allied health.
PMS vendor bundles recall	Incumbent adds AI features.	Stay ahead on outcome reporting, multilingual depth and integrations; partner rather than compete where possible.

Retention playbook

WEEK ONE

Deliver a visible win.

Import, audit and first approved campaign within the first week. Show the practice one booked appointment it can trace.

EVERY WEEK

Send the scorecard.

Response, bookings, recovered production and escalations. Make the number impossible to miss.

EVERY MONTH

Review the holdout.

Compare treated and control groups, explain the assumptions and agree on one expansion.

03 VALIDATION PLAN

Validate before writing production code.

The product plan assumes the business case is validated in parallel with the first build. The audit is both the lead magnet and the validation instrument.

Week	Action	Success criterion
1	30 interviews with practice managers in Germany and the Netherlands; ask for current recall process, overdue patient count, staff hours, PMS and willingness to pay.	At least 20 interviews completed; 10 practices share a real overdue-patient number.
2	Build a free “Recall Revenue Audit” using CSV upload. Run it for 5 clinics. Deliver a one-page report with estimated revenue at risk.	5 audits delivered; at least 3 clinics say the number is credible.
3	Pre-sell 10 pilots at €199/month for 90 days, 50% upfront. Use a concierge workflow: manual segmentation, AI-assisted messages, human review.	10 paying pilots, or 5 pilots plus 5 strong written commitments.
4	Run the first pilot messages, measure response and bookings, interview the front desk and document every objection.	20%+ response rate; 10+ booked appointments in the first month; front-desk time under 2 hours/week.

Kill criteria

Do not build the full product if fewer than 3 clinics will pay after 30 interviews; the response rate is below 10%; fewer than 5 appointments are recovered per clinic per month; or the PMS integration cannot be solved with CSV or one API within 6 weeks. In that case, pivot to veterinary clinics with the same engine or to SwitchboardAI.

Operating expenses and break-even

Stage	Monthly OPEX	What it buys
0–10 customers	€8k–15k	Founders, infrastructure, tools, legal/DPA, minimal marketing. No salaries beyond founders.
50 customers	€25k–35k	2 founders, 1 engineer, 1 part-time CS/support, tools, marketing experiments.
100 customers	€42k–58k	4 people, infrastructure, marketing, accounting/legal, co-working. MRR €34.9k is not enough.
250 customers	€70k–95k	7–8 people, partner manager, support, marketing. MRR €87.3k; roughly break-even.
500 customers	€110k–150k	10–12 people, sales, support, engineering, partnerships. MRR €174.5k; profitable if churn stays below 4%.

Planning estimates, not verified figures. Salaries vary significantly by country.

Resolve before billing

The landing page currently lists **5,000 patients/month** on the Scale plan; the business plan lists **6,000**. Pick one source of truth in the product’s plan catalog and update the landing page before Stripe billing goes live.

04 PRODUCT PLAN

The MVP is a closed loop, not a platform.

The first release must take a practice from a CSV export to a credible recovered-revenue report. Everything else is deferred until the loop is proven with real clinics.

MVP scope

Area	MVP behavior
Auth & tenancy	Org signup/login, invitations, owner/admin/staff roles, locations, providers.
Data import	CSV upload, column mapping, validation, consent fields, idempotent re-import.
Free audit	Overdue patients, at-risk production estimate, sample messages, shareable report link.
Recall engine	Hygiene overdue, treatment follow-up, lapsed 12–24 months, due in 30 days; per-provider/location rules.
Campaigns	Segment selection, holdout group, email templates, AI-assisted drafts, human approval, quiet hours, caps.
Inbox	Threaded conversations, intent classification, suggested replies, human takeover, escalation queue.
Booking	Public booking page with attribution token, confirmation email/ICS, staff confirmation.
Dashboard	Recovered appointments, recovered production, funnel, recall pipeline, agent status.
Reporting	Monthly PDF with treated vs control-group lift.
Billing	Stripe Checkout/Portal, plan limits, usage event tracking.
Compliance	Audit log, opt-out suppression, retention controls, RLS, DPA/DPIA artifacts.

OUT OF SCOPE FOR V0.1

- Voice calling and TTS
- WhatsApp and SMS (email-only first)
- Open Dental/Dentrix/Exact/SOE write-back
- Insurance verification, clinical advice, diagnosis
- Native mobile apps, enterprise SSO, multi-currency
- n8n for critical paths; agent frameworks

WHY EMAIL-FIRST

Email is the cheapest channel, needs no sender registration and is enough to prove response and booking lift. SMS and WhatsApp add deliverability, consent and template complexity that would slow the first pilot. Add them after the loop works.

04 RECALL WORKFLOWS

The recall loop, step by step.

Every step produces a record the practice can inspect. Nothing is sent without an eligibility check, and every booking can be traced back to a campaign and a message.

Core workflows

1 **Audit**

Upload CSV without full account → normalize and validate → overdue cohorts and at-risk production → sample messages → shareable report → convert to org.

2 **Campaign**

Rules engine creates tasks → campaign with 90/10 holdout → eligibility checks → LLM drafts within approved templates → human approval → email send → inbound reply → classification/escalation → booking.

3 **Booking**

Patient picks a slot or requests a time → booking record plus attribution token → confirmation/ICS → front desk confirms. PMS write-back comes later.

4 **Attribution**

Compare treatment vs holdout booking rates. Report gross attributed revenue and control-adjusted lift separately. Label small samples as directional.

Attribution formula

lift = treated booking rate − control booking rate. Incremental revenue = lift × treated patients × average production. Report gross attributed revenue separately and never present it as causal without the control group.

Before a message is sent

Check	Rule
Consent	Channel consent must be granted for the selected channel. Unknown or expired consent blocks the send.
Opt-out	Any prior opt-out for the patient or channel suppresses all future messages, regardless of campaign.
Quiet hours	Send only inside clinic-defined windows in the clinic timezone, including local holidays.
Frequency cap	Maximum messages per patient per 30 days; no duplicate outreach across active campaigns.
Recent appointment	Skip patients already booked or seen since the task was generated.
Open conversation	Do not start a new thread while a human is replying or an escalation is unresolved.

04 TECHNICAL ARCHITECTURE

One TypeScript codebase, few moving parts.

The brief proposes Next.js + NestJS + FastAPI + Redis + n8n. For a small team, simplify to a modular Next.js app, a Postgres-backed worker and shared domain packages. Add Python only when voice or heavier NLP requires it.

Layer	Choice	Rationale
Web	Next.js 15 App Router, TypeScript, Tailwind, shadcn/ui	SSR booking/audit pages, server actions, fast dashboard development.
Auth / DB / Storage	Supabase EU (Postgres, Auth, RLS, Storage)	EU residency, built-in auth, row-level tenant isolation.
ORM	Drizzle + zod	Type-safe schema and shared validation across web and worker.
Jobs	pg-boss on Railway	Scheduling, retries and rate limits without Redis at MVP.
AI	Vercel AI SDK; OpenAI primary, Claude fallback; structured outputs	Provider redundancy and cost negotiation. No fine-tuning.
Email	Resend or Postmark behind an adapter; inbound webhook	Start with email only; swap providers without touching domain logic.
Later channels	Twilio adapter for SMS/WhatsApp	Same channel interface as email.
Billing	Stripe Checkout + Customer Portal + webhooks	Subscriptions, trials, dunning, self-serve portal.
Reports	react-pdf; Recharts and TanStack Table	Monthly PDF and data-dense dashboard.
Observability	Sentry + PostHog EU	Errors, activation funnel and product analytics.
Deploy	Vercel (web), Railway (worker), Supabase migrations	Low ops, EU regions, preview deploys.

Why not the full stack from the brief?

NestJS, FastAPI, Redis and n8n add deployment and operational overhead before the product has proven demand. Next.js route handlers plus a pg-boss worker are enough for the first thousand patients. Revisit the split when voice, heavy document processing or throughput demand it.

04 DATA AND COMPLIANCE

Tenant isolation is the foundation.

Patient data is sensitive. The schema, RLS policies, retention jobs and audit trail are product requirements, not infrastructure details.

Product repository layout

```
RecallPilot-App/
├─ apps/
│  ├─ web/                # Next.js 15: dashboard, inbox, booking, audit, APIs
│  └─ worker/             # pg-boss: imports, sends, classification, reports
├─ packages/
│  ├─ db/                 # Drizzle schema, migrations, seeds, RLS tests
│  ├─ core/               # rules, segments, attribution, plans, i18n, zod
│  ├─ integrations/       # PMS and channel adapter interfaces
│  └─ ui/                 # shared shadcn components and brand tokens
├─ supabase/               # local config and policies
├─ .github/workflows/      # CI
└─ package.json            # npm workspaces
```

Minimum data model

- **Tenancy:** organizations, locations, users, memberships, providers
- **Patients:** patients, patient_channel_consents, appointments
- **Recall:** recall_rules, recall_tasks
- **Outreach:** campaigns, campaign_members, conversations, messages, message_events
- **Outcomes:** bookings, attribution_results, report_runs
- **Operations:** imports, import_rows, integration_connections, opt_outs, usage_events, subscriptions, plan_catalog, audit_logs

Attribution mechanics

- Randomize campaign members 90/10 into treatment/holdout.
- Stratify by location, treatment type and overdue bucket.
- Outcome window: 30 days initially, later 60/90.
- Revenue from PMS production value, or clinic-configured average.
- Booking tokens prevent double counting and tie revenue to the campaign.
- Show confidence and mark small samples as directional.

Security and compliance

- Supabase EU region; RLS on every tenant table; cross-tenant tests in CI.
- Integration credentials in Supabase Vault; service keys only on server/worker.
- Consent enforced per channel; unknown consent blocks sending; global opt-out suppression.
- Quiet hours in clinic timezone; frequency caps; SPF/DKIM/DMARC; List-Unsubscribe.
- Clinical, insurance, complaint and ambiguous replies always escalate to a human.
- Append-only audit log; PII redaction in Sentry/logs; retention and deletion jobs.
- DPA/DPIA templates, subprocessor list and incident runbook before scale.
- Raw audit-lead uploads purged on a short retention schedule.

05 DELIVERY PLAN

Pilot-ready in 10–12 weeks.

Assumes one full-time developer. Validation runs in parallel so the build is never more than one milestone ahead of real clinic feedback.

R0	**Bootstrap** Repo, workspaces, Next.js app, worker, Supabase EU, auth, CI, brand tokens. New user creates an org; RLS blocks cross-tenant access.	Week 1
R1	**Audit slice** CSV importer, normalized model, public audit report, share link, synthetic seed. A real clinic CSV produces a reproducible audit.	Weeks 2–3
R2	**Outreach slice** Rules engine, tasks, templates, AI drafts, approval queue, email adapter, worker. Scheduled campaign reaches a test inbox with approval and opt-out suppression.	Weeks 4–6
R3	**Inbox and booking** Reply webhooks, classification, escalation queue, patient timeline, booking page, ICS. Reply is classified, escalated and the booking is attributed.	Weeks 7–8
R4	**Dashboard and report** Recovery dashboard, holdout attribution, monthly PDF, audit log. Dashboard reflects real pilot data and the report shows treated vs control.	Weeks 9–10
R5	**Billing and hardening** Stripe, plan limits, usage metering, roles, e2e tests, security checklist. A pilot clinic can subscribe and operations are safe.	Weeks 11–12
R6	**Post-pilot expansion** Open Dental adapter/write-back, SMS/WhatsApp, multi-location, voice add-on. Only after pilot thresholds are met.	Months 4–6

Cross-repo integration

The landing repo remains unchanged. The product repo copies the logos, favicons, Geist fonts and color tokens once. Product URLs live at **app.recallpilot.com** with **/audit**, **/login**, **/signup** and **/book/[org]**. Later, landing CTAs move from **#get-started** to the product signup and audit URLs. The plan catalog in **packages/core/src/plans.ts** is the single source of truth for pricing and limits.

05 PILOT SCORECARD

What “working” looks like.

The pilot is successful only if the practice sees credible recovered revenue and the front desk does not absorb the work. These thresholds decide whether the next milestone is funded.

Release success criteria

Dimension	Target	Measured by
Audit credibility	At least 3 of 5 clinics say the overdue-patient and revenue-at-risk numbers are credible.	Audit interviews and conversion to pilot
Pilot demand	10 paying pilots at €199/month for 90 days, or 5 pilots plus 5 written commitments.	Signed pilot agreements and upfront payments
Patient response	20%+ response rate on the first pilot messages.	Conversation and reply events
Recovered appointments	10+ booked appointments per clinic in the first month.	Bookings tied to campaign tokens
Front-desk load	Under 2 hours per week of staff time during the pilot.	Inbox handling time and staff interviews

Working cadence

WEEKLY

One demo, one decision.

Review response, bookings, escalations and front-desk minutes. Choose one improvement for the next week.

MONTHLY

Recovered-revenue report.

Send the treated-versus-holdout report with assumptions and a clear next action.

GATE

Evidence before scope.

Do not start the next integration until the current milestone meets its threshold.

05 R0–R1 BUILD PLAN

Start with the audit.

The first milestone is the free audit plus CSV import. It is the lead magnet, the onboarding path and the data foundation for outreach.

R0 — Repo bootstrap

1. **Create the repo**
   Private v-cut127/RecallPilot-App; clone locally; switch the session to the new workspace.
2. **Set up npm workspaces**
   apps/web (Next.js 15), apps/worker (pg-boss), packages/db, core, integrations, ui.
3. **Copy brand assets**
   Logos, favicons, Geist fonts and color variables from the landing repo.
4. **Provision Supabase EU**
   EU Central (Frankfurt), local Supabase config, migrations and RLS policies.
5. **Auth and tenancy**
   Magic-link auth, organization creation, memberships and RLS tests.
6. **CI and environments**
   Typecheck, lint, tests, build; Vercel and Railway projects; .env.example.

R1 — Audit and CSV import

PUBLIC FLOW

- **/audit** — lead capture, CSV upload, progress state.
- **/audit/[token]** — shareable report, at-risk production, assumptions, sample messages, PDF download.
- Report delivered by email; conversion CTA to create an org.

AUTHENTICATED FLOW

- **/onboarding** — practice, locations, import, recall rules summary.
- **/app/import** — delimiter/encoding detection, column mapping, validation preview, commit.
- **/app/dashboard** — placeholder with import counts and audit link.

Import pipeline

1 · INGEST
Signed upload to Storage; file size/type limits; virus scan hook; raw file marked for deletion after processing.

2 · PARSE
UTF-8/BOM, comma/semicolon delimiters, European date formats, quoted fields, formula-injection protection.

3 · MAP
Auto-detect columns with manual override; save mapping profiles per PMS export.

4 · VALIDATE
Duplicates, missing contacts, invalid email/phone, missing consent, unmatched locations/providers.

5 · NORMALIZE
Write patients, consents and appointments; store import errors for download.

6 · AUDIT
Overdue by treatment type; lapsed 12–24 months; due next 30 days; production at risk with explicit assumptions.

05 R1 ACCEPTANCE

Done means a clinic can trust the number.

R1 is complete when a practice can upload its own data, see an audit it believes, and share it with a colleague without manual help from the team.

R1 acceptance criteria

- ✓ A 5,000-row clinic CSV imports and produces an audit in under 60 seconds.
- ✓ Mapping and validation errors are visible and downloadable; no silent data loss.
- ✓ RLS tests prove one org cannot access another org's patients or reports.
- ✓ Report token is unguessable, revocable and expires.
- ✓ The report is readable on mobile and labels every estimate as an assumption.
- ✓ A new user can go from signup to an imported audit without manual intervention.
- ✓ Synthetic seed covers 3 locations, ~2,500 patients, messy rows and two CSV variants.
- ✓ Raw audit-lead uploads are purged on the defined retention schedule.

Synthetic data for R1

Until a real clinic export is available, seed the product with a realistic synthetic clinic: 3 locations, ~2,500 patients, hygiene and treatment recall gaps, consent flags, duplicates and deliberately malformed rows. Export it in two variants — comma/ISO dates and semicolon/DD.MM.YYYY — so the importer is tested against European formats while the UI stays English.

Demo script for the first clinic

01
Upload the synthetic CSV and show the auto-detected columns.

02
Show validation: duplicates, missing consent, invalid contacts.

03
Open the audit: overdue patients and production at risk.

04
Share the tokenized report and create the org account.

05 RISKS AND NEXT STEPS

The risks are operational, not technical.

The stack is conventional. The hard parts are attribution credibility, deliverability, consent and front-desk adoption. Build only as far ahead as the evidence supports: the audit earns outreach, outreach earns booking, and pilots earn deeper integrations.

Risk	Why it matters	Mitigation
Attribution credibility	Weak numbers destroy the wedge and renewals.	Holdout group from day one; show assumptions; label directional results; booking tokens.
Email deliverability	Pilot fails before the agent gets a chance.	Dedicated sending domain, SPF/DKIM/DMARC, warm-up, bounces/complaints monitoring, provider adapter.
Consent and GDPR	One unlawful message can end the company.	Per-channel consent enforcement, suppression lists, DPA/DPIA, retention, audit logs, legal review.
Import variability	Every PMS exports different columns and dates.	Mapping profiles, validation reports, real CSV fixtures, normalized model, CSV fallback.
LLM safety	Clinical or insurance promises are unacceptable.	Approved templates, policy validator, clinical escalation, human approval in pilot mode, full logging.
Front-desk adoption	Staff revert to phone calls if the inbox adds work.	Human-in-the-loop queue, bulk approve, suggested replies, training, quick wins in week one.

Immediate next steps

1. **Create the product repo**
   v-cut127/RecallPilot-App, private, with npm workspaces.
2. **Provision the environment**
   Supabase EU, Stripe test mode, Resend, PostHog EU, Sentry.
3. **Ship R0**
   Auth, organizations, RLS, CI and the synthetic seed.
4. **Ship R1**
   CSV import and the free audit report, then put it in front of five clinics.
5. **Run validation in parallel**
   Interviews, audits and pilot pre-sales; enforce the kill criteria before building deeper integrations.

06 BRAND SYSTEM

Quiet confidence, measurable outcomes.

The brand is restrained and clinical without feeling cold: large Geist headlines, generous white space, ink-dark text, soft blue surfaces and a single confident accent. The product should feel like a reliable operating tool, not a chatbot.

Logo usage



Use the black logo on white and light surfaces. Minimum clear space: 1× the height of the mark.



Use the white logo on **#00061C** and dark imagery. Never place the black logo on a dark background.

Light palette

Background
#FFFFFF
Page background

Card
#F5F8FC
Panels, cards, tables

Pill
#EBF2F9
Tags and badges

Border
#DCE5EF
Hairlines and dividers

Muted
#526078
Secondary text

Ink
#00061C
Headlines and body

Accent
#006EC3
Links, labels, emphasis

Accent gradient
#006EC3 → #53D2E5
Hero atmosphere

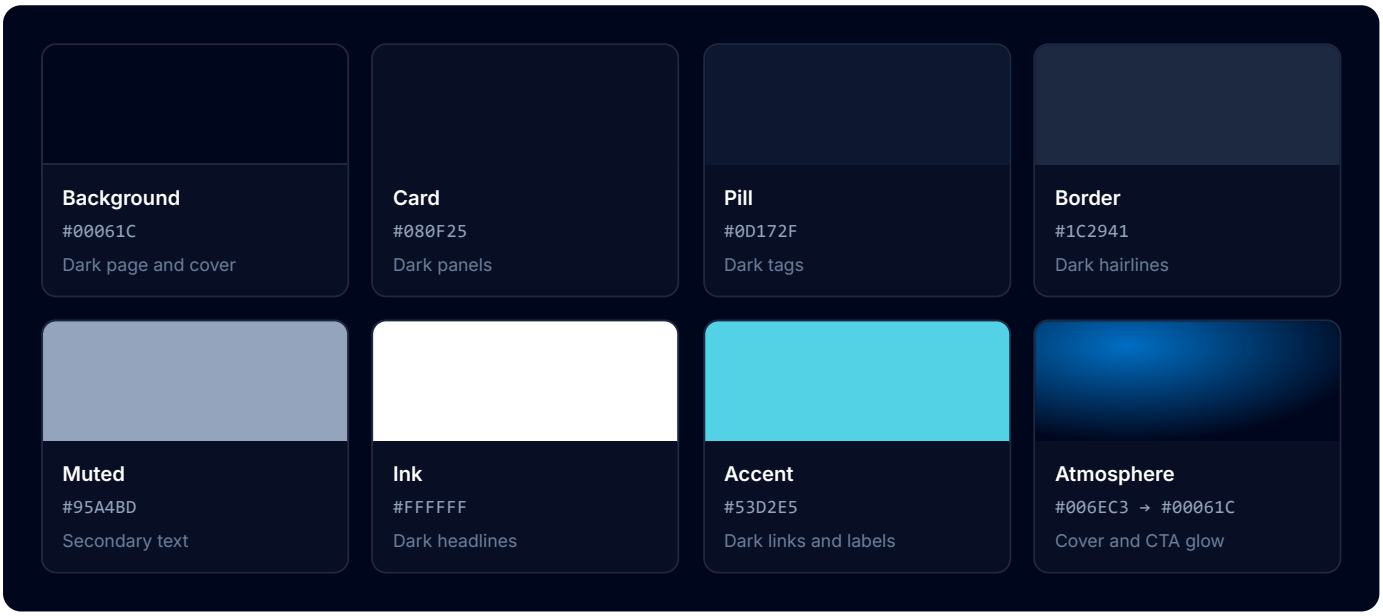
Usage rules

Never stretch the logo, recolor it, add effects or place it on a low-contrast background. Use the accent blue sparingly: links, active states, key labels and one primary CTA per view. Use ink for primary text and muted for secondary text; avoid introducing new grays.

06 DARK PALETTE

The dark system carries the cover and product shell.

Dark surfaces are used for the cover, the app shell and high-emphasis moments. Keep the hierarchy identical to the light theme: background, card, pill, border, muted, ink and accent — only the values change.



Contrast rules

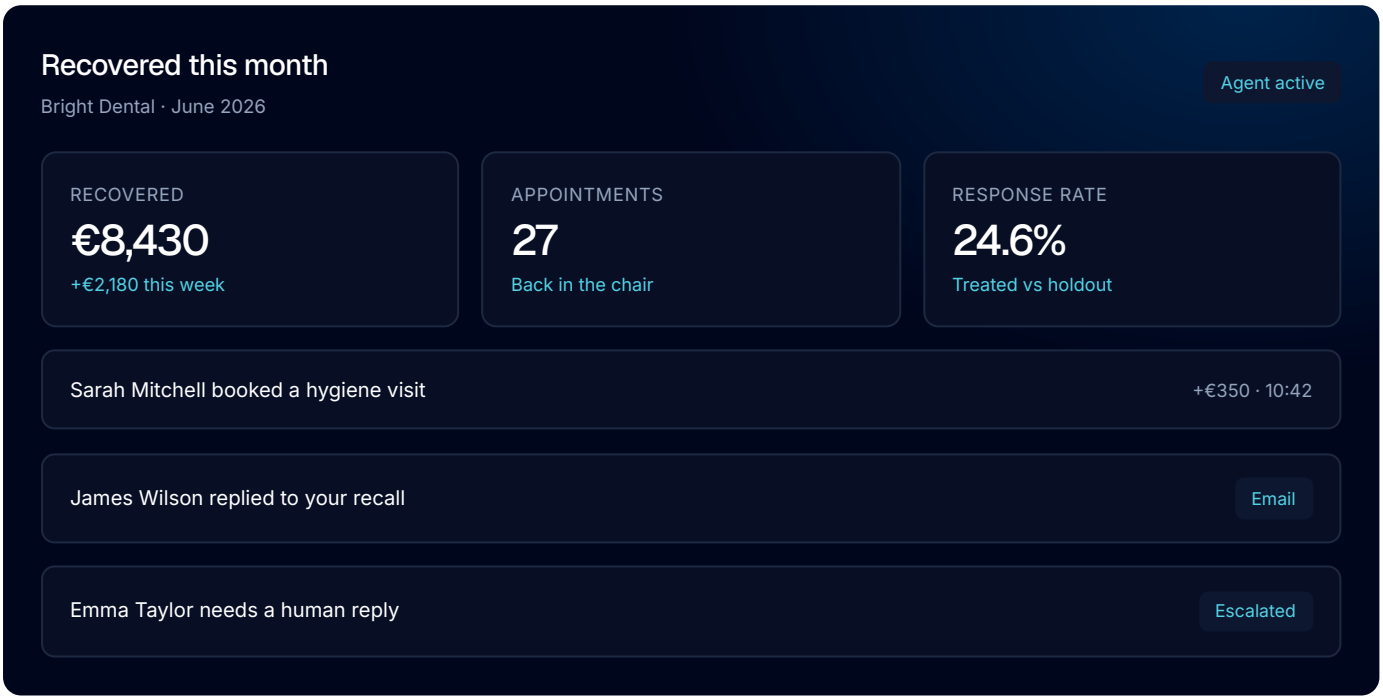
Use white or **#95A4BD** for text on **#00061C**. Reserve **#53D2E5** for links and active states; never use the light-theme accent **#006EC3** on dark surfaces. Keep body copy at 13px or larger in the product and 9pt or larger in print.

06 BRAND IN PRODUCT

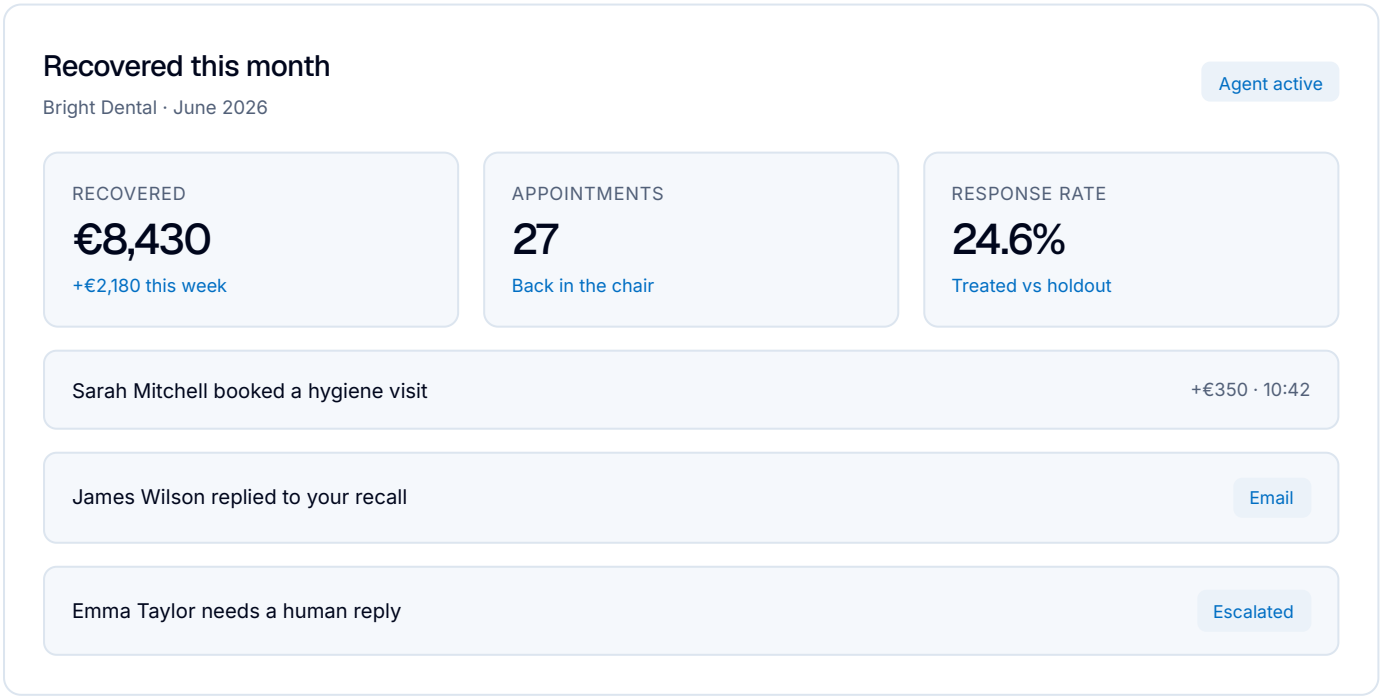
The same hierarchy in light and dark.

The product shell can run in either mode without changing the information hierarchy. Keep the cards, labels and accent behavior consistent; only the surface values change.

Dark workspace



Light workspace



06 TYPOGRAPHY AND UI TOKENS

Geist for decisions. Inter for detail.

Geist carries headlines, numbers and navigation. Inter carries body copy, tables and long-form content. Both are open-source and self-hosted with the product.

Geist

400 / 500 / 600 · woff2

Recovered this month

A plan for your practice. Bring patients back.
0123456789 €
ABCDEFGHIJKLMNOPQRSTUVWXYZ
abcdefghijklmnopqrstuvwxyz

Regular	Medium	Semibold
Headlines, large numbers	Default UI and headings	**Emphasis and labels**

Inter

400 / 600 / 700 · woff2

Recovered this month

Body copy, instructions, table cells and report prose. Inter keeps dense operational screens legible at small sizes.

Regular	Semibold	Bold
Body and tables	**UI labels and emphasis**	**Rare emphasis**

Type scale

Role	Font	Size	Usage
Display	Geist 500	44–68 px	Marketing hero and major statements
Section title	Geist 500	27–44 px	Page and section headings
Card title	Geist 500	22 px	Feature and plan cards
Metric	Geist 500	28–30 px	Dashboard KPIs and prices
Body	Inter 400	13–17 px	Product copy, tables, settings
Caption / label	Geist 400–600	11–13 px	Metadata, tags, table headers

Surfaces and components

Get started

Integration to be confirmed

Recall dashboard

Card
#F5F8FC · 1px #DCE5EF · 10px radius

10px

Card radius. Buttons use 7px; pills use 5px.

24px

Default card padding on desktop; 16px on compact cards.

1px

Hairline borders using the border token, never pure gray.

06 ASSETS AND VOICE

The words matter as much as the mark.

The brand is compliant, calm and specific. Product copy should help a practice manager understand what happened, what it is worth and what to do next.

Asset library

Asset	Location in the landing repo	Product use
Black logo	**public/assets/logo-black.png**	Light UI, documents
White logo	**public/assets/logo-white.png**	Dark UI, cover
Favicons	**public/favicon-light.png**, **favicon-dark.png**	Browser tabs
Touch icon	**public/touch-icon.png**	iOS home screen
Geist 400/500/600	**public/fonts/geist-*.woff2**	Headings and UI
Inter 400/600/700	**public/fonts/inter-*.woff2**	Body and tables

Voice and tone

CLEAR

Say the number.

Lead with recovered appointments and production. Avoid vague claims like “streamline” or “leverage AI.”

CALM

Never pressure the patient.

Messages are helpful, short and easy to ignore. Opt-out is always one click or one word away.

CAREFUL

Escalate, don’t guess.

Clinical, insurance and complaint language is never improvised. The agent hands off to a human.

Writing rules

Use sentence case. Prefer short sentences and concrete nouns. Show the assumption behind every estimate. Never imply clinical advice, guaranteed coverage or guaranteed revenue. Localize with native review before sending in German or Dutch.

Copy examples

Instead of	Say
“Leverage AI to streamline your recall.”	“Recover overdue appointments from your own patient list.”
“Guaranteed revenue growth.”	“Estimated recovered production, with the assumptions shown.”
“Your patients are overdue.”	“Three simple ways to get back on track.”
“Book now!!!”	“Choose a time that works for you.”

RecallPilot

Build the audit.
Earn the right to build the agent.

The first release is the wedge: a credible number that tells a practice how much production is waiting in its own patient database. Everything else follows from proving that number.